

Execution Blueprint

Sprint, Specs, Pricing, Rebrand

A comprehensive 12-week execution plan for transforming Datiq from a URL extraction tool into a full web and data intelligence platform, covering social listening, CRM connectors, and modular commercialization.

Prepared for: Datiq Product & Engineering

August 2026 | Version 1.0

1. Twelve-Week Sprint Plan: Phase 0 and Phase 1	3
1.1 Phase 0: Foundation and Infrastructure (Weeks 1-4)	3
1.1.1 Sprint 0.1 (Weeks 1-2): API Foundation and Rebrand Assets	3
1.1.2 Sprint 0.2 (Weeks 3-4): Data Pipeline, Billing, and Rebrand Launch	4
1.2 Phase 1: Core Intelligence Modules (Weeks 5-12)	5
1.2.1 Sprint 1.1 (Weeks 5-6): Social Listening MVP Core	6
1.2.2 Sprint 1.2 (Weeks 7-8): Connectors and Cross-Module Integration	7
1.2.3 Sprint 1.3 (Weeks 9-10): Pricing Launch and Social Enhancements	8
1.2.4 Sprint 1.4 (Weeks 11-12): Polish, Analytics, and Phase 2 Prep	9
1.3 Sprint Summary and Resource Allocation	10
2. Social Listening MVP: Feature Specification	11
2.1 Data Sources and Integration Architecture	11
2.2 Database Schema Design	12
2.3 User Interface Specification	13
2.3.1 Main Dashboard Layout	14
2.3.2 Mention Feed and Filtering	14
2.3.3 Alert Configuration Interface	14
3. Pricing and Packaging Model with Revenue Projections	15
3.1 Tier Structure and Feature Matrix	15
3.2 Revenue Projections and Unit Economics	16
3.3 Unit Economics Deep-Dive	17
4. Homepage Rebrand Copy and Visual Hero Concepts	18
4.1 Tagline Evaluation and Recommendation	18
4.2 Complete Homepage Copy	18
4.2.1 Hero Section Copy	19
4.2.2 Value Propositions Section	19
4.2.3 Social Proof and Trust Section	19
4.2.4 Module Overview Section	19

4.2.5 Pricing Section and Final CTA	20
4.3 Three Visual Hero Concepts	20
4.3.1 Concept A: "The Dashboard Reveal" (Recommended)	20
4.3.2 Concept B: "The Data Constellation"	21
4.3.3 Concept C: "The Before/After Split"	21
5. Connector Implementation Plan: HubSpot, Apollo, Stripe, and Calendly	22
5.1 Connector Architecture Overview	22
5.2 HubSpot Connector (Priority: Critical)	23
5.3 Apollo.io Connector (Priority: High)	24
5.4 Stripe Connector (Priority: Critical for Billing)	25
5.5 Calendly Connector (Priority: Medium, Phase 2)	26
5.6 Connector Implementation Timeline and Dependencies	27

1. Twelve-Week Sprint Plan: Phase 0 and Phase 1

This chapter presents a detailed, sprint-by-sprint execution plan covering the first twelve weeks of Datiq's transformation from a single-URL extraction tool into a modular web and data intelligence platform. The plan is divided into two phases: Phase 0 (Foundation and Infrastructure, Weeks 1-4) establishes the technical bedrock, API layer, and rebrand assets; Phase 1 (Core Intelligence Modules, Weeks 5-12) delivers the first revenue-generating features including Social Listening MVP, connector ecosystem, and updated pricing tiers. Each sprint is two weeks long, with clear deliverables, responsible owners, and engineering effort estimates calibrated for a small, focused team of 3-5 engineers plus one designer.

1.1 Phase 0: Foundation and Infrastructure (Weeks 1-4)

Phase 0 is the critical preparation phase. Before any user-facing intelligence features can ship, the platform needs a robust API layer, a modular architecture that supports multiple product modules, a rebranded identity system, and the database schema extensions required for social data, connectors, and billing. The goal is to emerge from Week 4 with a deployable REST API, a rebranded homepage, and the data infrastructure ready for Phase 1 feature development. Rushing past this phase to build features on fragile infrastructure is the single most common cause of startup platform rewrites.

1.1.1 Sprint 0.1 (Weeks 1-2): API Foundation and Rebrand Assets

The first sprint focuses on two parallel workstreams: building the public REST API layer and creating the rebrand design assets. The REST API is the single highest-leverage investment because it unlocks programmatic access, webhook integrations, Zapier/Make connectivity, and future mobile SDKs. Simultaneously, the design team produces the new brand kit (logo, color system, typography, component library) so that all subsequent feature UI is built on the new identity rather than the old one, avoiding a costly second re-skin later.

Table 1.1: Sprint 0.1 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P0-01	Design and implement REST API v1 endpoints for extraction (POST /api/v1/extract), batch (POST /api/v1/batch), and schedules (GET/POST /api/v1/schedules)	Backend Lead	5	None
P0-02	Implement API key management in Supabase (generate, revoke, rate-limit per tier)	Backend Lead	2	P0-01
P0-03	Build API documentation with OpenAPI 3.0 spec and interactive Swagger UI	Backend + Designer	2	P0-01
P0-04	Create new brand identity system: logo, primary/secondary palette, typography (Inter for UI, Playfair Display for marketing), icon set, and spacing tokens	Designer	4	None
P0-05	Build reusable Figma component library with new brand tokens for buttons, cards, tables, modals, and navigation	Designer	3	P0-04
P0-06	Set up modular routing architecture in Next.js app (datiq.app/extract, /social, /brand, /connect, /settings)	Frontend Lead	3	None
P0-07	Extend Supabase database schema: add tables for social_mentions, connectors, webhooks, billing_events, and audit_logs	Backend Lead	3	None
P0-08	Implement Redis caching layer for extraction results and API responses (TTL-based, per-user quotas)	Backend Lead	2	P0-01

The total engineering estimate for Sprint 0.1 is approximately 24 person-days. With a team of three engineers and one designer working in parallel on independent workstreams (API vs. design vs. frontend routing), this sprint is achievable within the two-week window, though the API documentation task may slip into Week 3 if API endpoint design encounters complexity around rate-limiting logic. The key risk is scope creep on the brand identity system: the designer should produce a functional, version-1.0 brand kit rather than pursuing pixel-perfection on every component variant.

1.1.2 Sprint 0.2 (Weeks 3-4): Data Pipeline, Billing, and Rebrand Launch

The second sprint of Phase 0 completes the infrastructure by building the data ingestion pipeline for social media sources, implementing the billing and subscription management system that will support the new pricing tiers (detailed in Chapter 3), and launching the rebranded homepage. The data pipeline is the most technically complex workstream because it must handle streaming data from multiple social platforms with different rate limits, authentication flows, and data formats. The billing system leverages Stripe's subscription API with usage-based metering for extraction credits and social mention quotas.

Table 1.2: Sprint 0.2 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P0-09	Build social data ingestion pipeline: Twitter/X API v2 connector, Reddit API connector, and generic RSS/Atom feed parser	Backend Lead	5	P0-07
P0-10	Implement real-time webhook system (user-configurable URLs triggered on extraction complete, mention detected, sentiment threshold breach)	Backend Lead	3	P0-01, P0-07
P0-11	Integrate Stripe Checkout and Customer Portal for subscription management (free, starter, pro, enterprise tiers)	Backend Lead	3	P0-07
P0-12	Build usage metering system: track extraction credits, social mention quotas, connector API calls per billing cycle	Backend Lead	2	P0-07, P0-11
P0-13	Implement new homepage with rebranded hero section, updated navigation, module overview cards, and pricing table	Frontend + Designer	4	P0-04, P0-05, P0-06
P0-14	Set up CI/CD pipeline with automated testing (unit tests for API, integration tests for data pipeline, E2E tests for key user flows)	Backend Lead	2	P0-01, P0-09
P0-15	Deploy rebranded homepage to production with feature flags for new modules (disabled until Phase 1)	Full Stack	1	P0-13
P0-16	Write internal API consumer documentation and integration guides for HubSpot and Apollo connectors (Chapter 5 spec)	Tech Writer	2	P0-03

Sprint 0.2 totals approximately 22 person-days. The critical path runs through P0-09 (social data pipeline) and P0-11 (Stripe integration), both of which have external dependencies on third-party API access and approval processes. The team should request Twitter/X developer access and Stripe sandbox credentials during Sprint 0.1 to avoid blockers. By the end of Phase 0, the platform will have a publicly documented REST API, a rebranded web presence, a billing engine ready for tiered monetization, and the foundational data schema to support all Phase 1 features.

1.2 Phase 1: Core Intelligence Modules (Weeks 5-12)

Phase 1 is where Datiq begins its transformation from a URL extraction utility into a multi-module intelligence platform. This phase delivers three major capabilities: the Social Listening MVP (detailed feature spec in Chapter 2), the first two CRM connectors (HubSpot and Apollo, implementation plan in Chapter 5), and the Stripe billing integration for the new pricing tiers (detailed in Chapter 3). Each of the

four sprints in Phase 1 is designed to produce a shippable increment that can be released to users, gathering feedback that shapes subsequent sprints. The sprint cadence follows a disciplined two-week cycle with sprint planning on Monday, daily standups, Thursday demos, and Friday retrospectives.

1.2.1 Sprint 1.1 (Weeks 5-6): Social Listening MVP Core

This sprint delivers the minimum viable version of Datiq Social, the platform's social media listening module. The MVP focuses on three data sources (Twitter/X, Reddit, and RSS feeds), keyword-based mention tracking, basic sentiment classification (positive, negative, neutral), and a real-time mention feed with filtering. This is intentionally narrow in scope to validate the core value proposition before investing in additional data sources or advanced analytics. The MVP must prove that users will engage with social data within the Datiq interface rather than switching to a dedicated tool like Brandwatch or Sprout Social.

Table 1.3: Sprint 1.1 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P1-01	Build Social Listening dashboard UI: mention feed, sentiment summary card, volume trend sparkline, keyword tag management panel	Frontend Lead	5	P0-06, P0-07
P1-02	Implement keyword/brand tracking engine: users add keywords, system matches against ingested social data, stores matched mentions	Backend Lead	4	P0-09
P1-03	Integrate sentiment analysis model (fine-tuned BERT or GPT-based classifier) into the ingestion pipeline for real-time classification	ML Engineer	4	P0-09, P1-02
P1-04	Build mention detail view: expanded post content, author info, source link, sentiment score, and related mentions	Frontend Lead	2	P1-01, P1-02
P1-05	Implement alert system: email and in-app notifications when mention volume spikes or negative sentiment exceeds threshold	Backend Lead	3	P1-02, P0-10
P1-06	Create Social Listening onboarding flow: guided setup wizard for first keyword, data source selection, and alert configuration	Frontend + Designer	3	P1-01
P1-07	Add social mention data to existing Dashboard module (cross-module aggregation with extraction stats)	Full Stack	2	P1-01, P0-08

Sprint 1.1 represents the highest-risk sprint in the plan because it introduces machine learning (sentiment analysis) alongside significant frontend and backend work. The sentiment model should be pre-trained during Phase 0 using publicly available Twitter and Reddit datasets so that Sprint 1.1 focuses on integration rather than model development. If the ML engineer is not available, a pragmatic fallback is to use OpenAI's text classification API as a temporary backend until a custom model is trained. The total estimate is 23

person-days, requiring full team capacity.

1.2.2 Sprint 1.2 (Weeks 7-8): Connectors and Cross-Module Integration

With the Social Listening MVP deployed, Sprint 1.2 shifts focus to the connector ecosystem. The two highest-priority connectors are HubSpot and Apollo.io, chosen because they represent the two most common workflows in Datiq's target market (SMB sales and marketing teams). The HubSpot connector enables bi-directional sync of extracted contact data and social intelligence into HubSpot CRM contacts and deals. The Apollo connector enriches Datiq extraction results with Apollo's 275-million-contact database, adding firmographic data, verified emails, and direct dial numbers to extracted leads. This sprint also integrates social intelligence data into the existing Collections module, allowing users to save and organize social mentions alongside URL extractions.

Table 1.4: Sprint 1.2 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P1-08	Build HubSpot OAuth 2.0 integration: authorization flow, token management, and bi-directional data sync for contacts and companies	Backend Lead	5	P0-01, P0-07
P1-09	Implement Apollo.io API integration: contact enrichment, company lookup, and sequence sync with extraction results	Backend Lead	4	P0-01
P1-10	Build connector management UI: authorization status dashboard, sync history log, field mapping configuration, and disconnect/reconnect controls	Frontend Lead	4	P1-08, P1-09
P1-11	Implement data mapping engine: transform Datiq extraction schema to HubSpot/Apollo field formats with user-customizable field maps	Backend Lead	3	P1-08, P1-09
P1-12	Extend Collections module to support social mention items alongside extraction results (unified item type with polymorphic storage)	Full Stack	3	P1-01, P0-07
P1-13	Build Stripe integration for Calendly-style usage metering: track social mentions consumed, connector API calls, and extraction credits per billing cycle	Backend Lead	3	P0-11, P0-12
P1-14	Implement automated data sync scheduling: users configure periodic sync intervals (hourly, daily, weekly) for each connector	Backend Lead	2	P1-08, P1-09

Sprint 1.2 totals approximately 24 person-days and is heavily backend-weighted. The HubSpot and Apollo integrations each require navigating their respective OAuth flows, rate limits, and data models. A key design decision is whether to implement a generic connector abstraction layer (additional 3-4 days upfront) or build each connector as a standalone module. The recommendation is to build a lightweight abstraction with a

standardized interface (authorize, sync, map_fields, disconnect) that both connectors implement, even if the internal logic differs. This investment pays dividends when adding Stripe, Calendly, and future connectors in subsequent phases. External dependencies include HubSpot App Marketplace approval (can take 1-2 weeks) and Apollo API access tier selection.

1.2.3 Sprint 1.3 (Weeks 9-10): Pricing Launch and Social Enhancements

Sprint 1.3 activates the monetization engine by launching the new pricing tiers (detailed in Chapter 3) and enhancing the Social Listening MVP based on user feedback from the Sprint 1.1 release. Key pricing tasks include building the self-service upgrade flow, implementing the tiered feature gates (which modules are accessible at each tier), and deploying the usage dashboard where users can monitor their credit consumption. On the social side, this sprint adds historical data backfill (users can import up to 30 days of historical mentions for their tracked keywords), competitive tracking (monitor competitor brand keywords alongside your own), and an improved sentiment visualization with trend comparison charts.

Table 1.5: Sprint 1.3 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P1-15	Build self-service pricing page with tier comparison, feature matrix, and Stripe Checkout integration for plan upgrades/downgrades	Frontend + Backend	4	P0-11, P0-12, P0-13
P1-16	Implement tiered feature gates in the application router: check user subscription tier before granting access to Social, Connectors, API, and Batch modules	Full Stack	3	P0-11, P0-06
P1-17	Build usage dashboard: visual breakdown of extraction credits consumed, social mentions processed, connector API calls made, and remaining quotas	Frontend Lead	3	P0-12, P1-13
P1-18	Add historical data backfill for social mentions: bulk import up to 30 days of historical mentions for tracked keywords from Twitter and Reddit	Backend Lead	3	P1-02, P0-09
P1-19	Implement competitive tracking: users add competitor brand names as tracked keywords with separate sentiment and volume analytics	Full Stack	3	P1-02, P1-03
P1-20	Build sentiment trend comparison charts: side-by-side trend lines for own brand vs. competitors with configurable date ranges	Frontend Lead	3	P1-18, P1-19
P1-21	Implement email digest reports: daily or weekly summary of social mentions, sentiment shifts, and volume anomalies delivered to user inbox	Backend Lead	2	P1-05, P1-18

Sprint 1.3 totals approximately 21 person-days and represents the revenue inflection point of the entire 12-week plan. For the first time, Datiq will have a clear, self-service path from free user to paying customer.

The pricing page and feature gates must be pixel-perfect and frictionless because every extra click in the upgrade flow directly reduces conversion. The competitive tracking feature is strategically important because it transforms Datiq from a single-brand tool into a competitive intelligence platform, positioning it against Klue (\$15-20K/year) and Crayon (\$25K+/year) at a fraction of the cost. This is the sprint where the platform narrative shifts from "extract data from URLs" to "understand your competitive landscape." The email digest feature is a retention mechanism: users who receive daily insights in their inbox are significantly less likely to churn.

1.2.4 Sprint 1.4 (Weeks 11-12): Polish, Analytics, and Phase 2 Prep

The final sprint of Phase 1 focuses on three objectives: polishing the user experience based on two sprints of feedback, implementing product analytics to measure engagement and inform Phase 2 prioritization, and preparing the technical foundation for Phase 2 features (brand positioning engine, social content writer, additional connectors). Polish work includes performance optimization for the social mention feed (lazy loading, virtual scrolling), accessibility improvements (WCAG 2.1 AA compliance for all new modules), and mobile responsiveness for the Social Listening and Connector management views. The analytics implementation tracks core SaaS metrics: DAU/MAU, feature adoption rates, time-to-value, and conversion funnel drop-off points.

Table 1.6: Sprint 1.4 Task Breakdown

Task ID	Task Description	Owner	Est. Days	Dependencies
P1-22	Implement product analytics with PostHog or Mixpanel: track DAU/MAU, feature adoption, conversion funnel, time-to-value, and retention cohorts	Backend Lead	3	P0-07
P1-23	Performance optimization: implement virtual scrolling for mention feed, lazy loading for extraction results, and Redis caching for API responses	Frontend + Backend	3	P1-01, P0-08
P1-24	Accessibility audit and remediation: ensure all new modules meet WCAG 2.1 AA (keyboard navigation, screen reader support, color contrast)	Frontend + Designer	3	P1-01, P1-10
P1-25	Mobile responsiveness: adapt Social Listening dashboard, connector management, and pricing page for tablet and mobile viewports	Frontend + Designer	3	P1-01, P1-10, P1-15
P1-26	Build internal admin dashboard: user management, subscription overview, system health monitoring, and usage analytics aggregation	Backend Lead	3	P0-11, P1-22
P1-27	Design Phase 2 technical architecture: brand positioning engine data model, social content writer prompt system, and connector abstraction layer v2	Tech Lead	2	All P1
P1-28	Conduct user research interviews (5-8 users) on Social Listening MVP and connector experience, document findings and Phase 2 recommendations	Product Manager	2	P1-01, P1-08

Sprint 1.4 totals approximately 19 person-days, the lightest sprint in the plan, intentionally so. After three intensive sprints of feature delivery, the team needs breathing room to address technical debt, fix bugs reported by early users, and gather the qualitative feedback that will shape Phase 2. The user research interviews are particularly valuable: they will validate or invalidate the assumption that SMB users want social listening and CRM integration within the same platform as web extraction, rather than using separate point solutions. The Phase 2 architecture design ensures the team can hit the ground running in Week 13 without a planning gap.

1.3 Sprint Summary and Resource Allocation

The following table summarizes all six sprints across Phase 0 and Phase 1, providing a bird's-eye view of effort distribution, key deliverables, and cumulative progress. The total engineering investment for the 12-week plan is approximately 133 person-days, which translates to roughly 6.6 person-months of effort.

With a team of three full-time engineers and one part-time designer, this plan requires approximately 20 weeks of calendar time if running a single stream, but the parallel workstreams compress this to 12 weeks. The plan assumes no major scope changes or external blockers; a 15-20% buffer is recommended for each sprint to accommodate unexpected complexity.

Table 1.7: Sprint Summary Overview

Sprint	Weeks	Phase	Est. Person-Days	Key Deliverables
0.1	1-2	Phase 0	24	REST API v1, brand identity, modular routing, schema extensions
0.2	3-4	Phase 0	22	Social data pipeline, Stripe billing, rebranded homepage launch
1.1	5-6	Phase 1	23	Social Listening MVP: mention feed, sentiment, alerts
1.2	7-8	Phase 1	24	HubSpot + Apollo connectors, connector management UI
1.3	9-10	Phase 1	21	Pricing launch, competitive tracking, email digests
1.4	11-12	Phase 1	19	Analytics, polish, mobile, Phase 2 architecture

2. Social Listening MVP: Feature Specification

This chapter defines the complete feature specification for Datiq Social, the social media listening module that represents the platform's first major expansion beyond URL extraction. The spec covers data sources and their integration methods, the database schema required to store and query social mentions, the user interface layout and interaction patterns, and the alert and notification system. The design philosophy is to start with three high-signal data sources (Twitter/X, Reddit, and web RSS) and deliver a polished experience on those before expanding to additional platforms. This approach mirrors how Brandwatch and Sprout Social started with Twitter data before adding Facebook, Instagram, and YouTube, except Datiq targets the SMB market with dramatically simpler setup and pricing.

2.1 Data Sources and Integration Architecture

The Social Listening MVP ingests data from three sources, each chosen for its relevance to the SMB target market, API accessibility, and signal quality. Twitter/X provides real-time brand conversation data with rich metadata (retweets, likes, replies, impressions). Reddit offers high-quality, long-form discussion threads that are particularly valuable for product feedback and sentiment analysis. RSS/Atom feeds cover news outlets, blogs, and industry publications that mention the tracked brand or keywords, filling the gap left by

social-only monitoring. Together, these three sources provide a comprehensive view of brand perception across social conversation and public media.

Table 2.1: Social Listening Data Sources

Data Source	API / Method	Rate Limit	Data Retention	Cost at Scale
Twitter/X	API v2 (Academic Research or Basic tier)	10,000 tweets/month (Basic), 1M (Academic)	30 days rolling window in MVP	\$100/month (Basic), \$5,000 (Academic)
Reddit	Reddit API (OAuth 2.0 via PRAW)	100 requests/minute	90 days rolling window	Free (no API cost at scale)
RSS/Atom Feeds	Server-side feed parser (rss-parser npm)	No rate limit (HTTP polling)	Unlimited (stored in DB)	Free (infrastructure only)
Future: LinkedIn	LinkedIn Marketing API (partnership required)	TBD (requires approval)	TBD	Requires LinkedIn Partnership Program
Future: News API	NewsAPI.org or GDELT Project	100 requests/day (free tier)	30 days	\$449/month (Pro tier)

The integration architecture follows an event-driven pattern. Each data source has a dedicated collector service that runs on a scheduled interval (Twitter: every 5 minutes, Reddit: every 15 minutes, RSS: every 30 minutes). Collectors fetch new mentions, pass them through a normalization layer that standardizes the data format, then enqueue them for sentiment analysis and storage. The normalization layer maps each source's unique schema to a unified mention object with the following fields: `source_platform`, `source_id`, `author_handle`, `author_name`, `content_text`, `content_html`, `published_at`, `url`, `engagement_metrics` (likes, shares, comments), and `extracted_keywords`. This abstraction allows new sources to be added by implementing a single collector interface without modifying downstream processing logic.

2.2 Database Schema Design

The social listening data model extends the existing Datiq Supabase schema with five new tables and one enum type. The design prioritizes query performance for the most common access patterns (filtering by keyword, date range, and sentiment) and supports the future addition of new data sources without schema migrations. All tables include the standard Datiq audit columns (`created_at`, `updated_at`) and reference the existing users and workspaces tables for multi-tenancy.

Table 2.2: Social Listening Database Schema

Table	Key Columns	Index	Relationship
tracked_keywords	id, workspace_id, keyword, label, is_competitor, source_filters (JSONB), created_by, is_active	UNIQUE(workspace_id, keyword), idx_source_filters (GIN)	Belongs to workspaces, has many social_mentions
social_mentions	id, workspace_id, keyword_id, source_platform, source_id, author_handle, author_name, content_text, sentiment (enum), sentiment_score (float), published_at, url, engagement_data (JSONB)	idx_workspace_keyword (workspace_id, keyword_id, published_at), idx_sentiment (sentiment), idx_platform (source_platform)	Belongs to workspaces and tracked_keywords
sentiment_trends	id, workspace_id, keyword_id, date, positive_count, negative_count, neutral_count, avg_score, total_mentions	UNIQUE(workspace_id, keyword_id, date)	Aggregation table, rollup from social_mentions daily
alert_rules	id, workspace_id, keyword_id, alert_type (spike/negative_threshold/volume), threshold_value, notification_channel (email/in-app/webhook), is_active	idx_workspace_active (workspace_id, is_active)	Belongs to workspaces and tracked_keywords
alert_events	id, alert_rule_id, workspace_id, triggered_at, mention_count, sentiment_value, notification_status, delivered_at	idx_rule_time (alert_rule_id, triggered_at)	Belongs to alert_rules

The sentiment enum type has three values: positive, negative, and neutral. The sentiment_score column stores a float between -1.0 (maximally negative) and +1.0 (maximally positive), with 0.0 representing neutral. This dual representation allows both categorical filtering (show only negative mentions) and quantitative analysis (calculate average sentiment trend over time). The engagement_data JSONB column stores platform-specific metrics (Twitter: retweet_count, like_count, reply_count; Reddit: upvotes, downvotes, comment_count; RSS: none) without requiring separate columns for each platform. The sentiment_trends table is a pre-computed aggregation table that stores daily rollups, avoiding expensive COUNT and AVG queries on the social_mentions table for dashboard rendering.

2.3 User Interface Specification

The Social Listening UI is accessed via a new top-level navigation item labeled "Social" in the Datiq app. The interface follows a three-panel layout: a left sidebar for keyword and filter management, a central feed for mention browsing, and a right panel for detail views and analytics. This layout is inspired by Sprout Social's Smart Inbox but simplified for the SMB user who does not need the multi-account management complexity of enterprise tools. The design uses the new Datiq brand tokens (from Sprint 0.1) and integrates visually with the existing extraction modules through shared card components, table styles, and the unified

navigation bar.

2.3.1 Main Dashboard Layout

The main Social Listening dashboard presents four key information zones. At the top, a summary bar displays total mentions in the last 7 days, average sentiment score (with a color-coded indicator), and the percentage change from the previous period. Below the summary, a volume trend sparkline chart shows mention volume over the last 30 days, with color-coded segments for positive (green), negative (red), and neutral (gray) mentions. The central area is the mention feed, a vertically scrollable list of social mention cards, each showing the source platform icon, author name and handle, content preview (truncated to 3 lines with expand), published time (relative: "2h ago", "1d ago"), sentiment badge, and engagement metrics. The right panel shows a sentiment distribution donut chart (positive/negative/neutral percentages) and a top keywords table ranking tracked keywords by mention volume. Users can click any keyword to filter the feed. The entire dashboard is responsive: on mobile, the panels stack vertically with the feed taking full width and analytics accessible via a bottom-sheet modal.

2.3.2 Mention Feed and Filtering

The mention feed supports real-time filtering across five dimensions: keyword (select one or more tracked keywords), sentiment (positive/negative/neutral/all), source platform (Twitter/Reddit/RSS/all), date range (preset options: last 24h, 7d, 30d, custom), and engagement threshold (minimum likes/shares/replies). Filters combine with AND logic: selecting "negative" sentiment and "Twitter" source shows only negative tweets. The feed supports two sort orders: newest first (default) and most engaging first (sorted by total engagement metrics). Each mention card is clickable, expanding to a detail view that shows the full content, author profile summary, direct link to the original post, sentiment score breakdown (which words/phrases contributed to the classification), and a "Save to Collection" button that adds the mention to an existing Datiq Collection alongside extraction results. The feed implements infinite scroll with a batch size of 20 mentions, loading the next batch when the user scrolls within 200px of the bottom. A "Load More" button appears after 100 mentions to prevent excessive DOM nodes.

2.3.3 Alert Configuration Interface

The alert configuration interface is accessible from the Social Listening settings panel, reachable via a gear icon in the dashboard header. Users create alert rules by selecting a tracked keyword, choosing an alert type (mention volume spike, negative sentiment threshold, or competitor mention), setting a threshold value, and choosing a notification channel. For volume spike alerts, the threshold is a percentage increase over the rolling 7-day average (default: 200%, meaning double the normal volume). For negative sentiment threshold alerts, the user sets a minimum negative percentage (default: 30% of mentions are negative). For competitor mention alerts, the system triggers when a tracked competitor keyword appears in the same conversation thread as the user's brand keyword. Each alert rule has an active/inactive toggle and a cooldown period (default: 1 hour) to prevent alert fatigue from sustained spikes. The notification channels include in-app notification (bell icon with badge count), email delivery (to the user's registered email), and webhook (POST to a user-configured URL with a JSON payload containing the alert details and affected mentions). Users can create up to 10 alert rules on the Starter tier and 50 on the Pro tier.

3. Pricing and Packaging Model with Revenue Projections

This chapter defines Datiq's monetization strategy for the expanded platform. The pricing model is designed around three principles: a generous free tier that drives adoption and word-of-mouth growth, a mid-tier that captures SMB teams willing to pay for social intelligence and CRM integration, and an enterprise tier for organizations that need advanced features, higher limits, and dedicated support. The model draws competitive pricing intelligence from across the market: Brandwatch starts at \$800/month, Sprinklr's listening module anchors around \$80,000/year, Klue and Crayon command \$15-25K/year, while Apollo.io's paid plans start at \$49/month. Datiq's positioning in the \$29-\$149/month range fills a massive gap between free tools and enterprise-only platforms, making web intelligence accessible to the 20 million SMBs worldwide that currently cannot justify five-figure social listening budgets.

3.1 Tier Structure and Feature Matrix

Table 3.1: Datiq Pricing Tier Feature Matrix

Feature / Limit	Free	Starter (\$29/mo)	Pro (\$79/mo)	Enterprise (Custom)
Monthly extractions	10	200	1,000	Unlimited
Batch extractions/month	5 runs	50 runs	200 runs	Unlimited
Scheduled extractions	1 schedule	10 schedules	50 schedules	Unlimited
Social mentions/month	Not included	2,000	10,000	Custom
Tracked keywords	Not included	5	25	Unlimited
Sentiment analysis	Not included	Basic	Advanced + trends	Custom models
Data sources (social)	Not included	Twitter + Reddit	All 5 sources	Custom sources
CRM connectors	Not included	1 connector (HubSpot or Apollo)	All connectors	Custom integrations
API access	Not included	REST API (1,000 calls/mo)	REST API (10,000 calls/mo)	Unlimited + webhooks
Team members	1	3	10	Unlimited + SSO
Data retention	30 days	90 days	1 year	Unlimited
AI training on data	Never	Never	Never	Never
Support	Community	Email (48h SLA)	Priority (4h SLA)	Dedicated CSM + Slack

The free tier is deliberately generous with URL extraction (10/month) to maintain the existing value proposition that drives sign-ups, but excludes all social intelligence and connector features. This creates a clear upgrade path: users who find value in extraction naturally want to extend that intelligence to social media and CRM, which requires moving to Starter. The Starter tier at \$29/month is priced below the psychological threshold of \$50, making it an impulse purchase for solo entrepreneurs and small teams. The Pro tier at \$79/month targets growing teams that need multiple connectors, higher limits, and advanced sentiment features. The Enterprise tier is custom-priced with annual contracts, targeting organizations that need SSO, custom data sources, dedicated support, and compliance features (SOC 2, data residency).

3.2 Revenue Projections and Unit Economics

Revenue projections are built on a bottom-up model that starts with user acquisition assumptions and applies conversion rates through the pricing funnel. The model assumes a phased launch: Phase 1 (months 1-6) focuses on the free tier and conversion optimization; Phase 2 (months 7-12) introduces paid tiers with marketing investment; Phase 3 (months 13-24) scales with paid acquisition and expansion revenue. The unit economics are calibrated against industry benchmarks: typical SaaS free-to-paid conversion rates range from 2-5%, monthly churn for SMB tools averages 3-5%, and net revenue retention for product-led growth companies targets 110-120%.

Table 3.2: Revenue Projections (24-Month Horizon)

Metric	Month 3	Month 6	Month 12	Month 18	Month 24
Total registered users	1,500	5,000	15,000	40,000	80,000
Active users (MAU)	600	2,000	7,500	20,000	40,000
Starter subscribers	15	75	300	800	1,600
Pro subscribers	3	20	100	350	750
Enterprise accounts	0	1	5	15	30
Monthly recurring revenue (MRR)	\$492	\$3,575	\$20,695	\$63,950	\$135,400
Annual run rate (ARR)	\$5,900	\$42,900	\$248,340	\$767,400	\$1,624,800
Free-to-paid conversion	3.0%	3.8%	4.0%	4.5%	4.8%
Monthly churn rate	5.0%	4.5%	4.0%	3.5%	3.0%
Customer acquisition cost (CAC)	\$45	\$35	\$28	\$22	\$18
Lifetime value (LTV)	\$87	\$156	\$237	\$316	\$395
LTV:CAC ratio	1.9x	4.5x	8.5x	14.4x	21.9x

The model projects Datiq reaching \$1.6M ARR by month 24, driven primarily by the Pro tier (\$79/month) which accounts for approximately 55% of revenue by month 24. The LTV:CAC ratio improves dramatically over time as organic word-of-mouth reduces acquisition costs and the multi-module platform increases switching costs (users who rely on Datiq for extraction, social listening, and CRM integration are far less likely to churn than single-feature users). The key assumption is that the social listening and connector features increase the free-to-paid conversion rate from the current estimated 1-2% (extraction only) to 4-5% (multi-module value proposition), which is consistent with industry data showing that multi-product platforms convert 2-3x better than single-product tools.

3.3 Unit Economics Deep-Dive

Understanding the per-customer economics is critical for validating the pricing model's sustainability. The cost structure for Datiq's expanded platform consists of four major components: infrastructure costs (Supabase hosting, Redis cache, API gateway), third-party API costs (Twitter/X, sentiment model inference, HubSpot/Apollo API calls), customer support costs, and sales and marketing expenses. Infrastructure scales linearly with active users, while third-party API costs scale with usage volume (social mentions ingested, connector syncs performed). The following analysis breaks down the cost-per-customer at each pricing tier.

Table 3.3: Unit Economics by Pricing Tier

Cost Component	Free User (Monthly)	Starter (\$29)	Pro (\$79)	Enterprise (est. \$499)
Infrastructure (Supabase + Redis)	\$0.50	\$0.80	\$1.50	\$8.00
Social data APIs (Twitter + Reddit)	\$0	\$3.20	\$12.00	\$45.00
Sentiment analysis (inference cost)	\$0	\$1.00	\$4.50	\$20.00
CRM connector API costs	\$0	\$1.50	\$5.00	\$15.00
Support cost (allocated)	\$0	\$2.00	\$5.00	\$50.00 (CSM)
Total cost per customer	\$0.50	\$8.50	\$28.00	\$138.00
Gross margin	N/A	71%	65%	72%

The gross margins of 65-72% are healthy for a SaaS platform at this stage and well above the 60% threshold that venture investors typically look for in SaaS business models. The Starter tier has the highest margin (71%) because its usage limits keep third-party API costs contained while the \$29 price point provides substantial margin. The Pro tier's margin is slightly lower (65%) due to higher social mention volumes and multiple connector costs, but the higher absolute dollar contribution ($\$79 - \$28 = \$51$ per customer per month) makes it the most profitable tier on a per-customer basis. The Enterprise tier's margin is the highest (72%) because enterprise pricing includes a significant premium for dedicated support, SSO, and compliance features that have relatively low marginal cost. The blended gross margin across all paying customers is approximately 68%, which supports sustainable growth reinvestment.

4. Homepage Rebrand Copy and Visual Hero Concepts

This chapter presents the complete rebrand copy for Datiq's homepage, including the new tagline, hero section copy, value proposition statements, and three distinct visual hero concepts with layout specifications. The rebrand shifts Datiq's positioning from "Intelligence from every URL" (which anchors the brand to a single, limited interaction pattern) to a broader platform identity that communicates web-scale intelligence, multi-source data synthesis, and actionable business insights. The new tagline candidates were evaluated against four criteria: memorability, brevity (under 6 words), differentiation from competitors, and ability to encompass all current and planned modules (extraction, social, brand, connectors).

4.1 Tagline Evaluation and Recommendation

After evaluating eight tagline candidates against the four criteria, the recommended primary tagline is **"Web Intelligence, Delivered."** This tagline scores highest on memorability (two-word punch), differentiation (no competitor uses "delivered" in this context), and encompassment ("web intelligence" covers extraction, social listening, brand monitoring, and any future module). The secondary tagline for marketing campaigns and SEO metadata is **"Understand the Web. Act on It."** which adds a call-to-action dimension that works well in paid advertising and email subject lines. The following table shows the full evaluation matrix.

Table 4.1: Tagline Evaluation Matrix

Tagline	Memorability	Brevity	Differentiation	Encompassment	Score (1-5)
"Web Intelligence, Delivered."	5 - Punchy, two-part rhythm	5 - 4 words	4 - "Delivered" is unique in category	5 - Covers all modules	4.8
"Understand the Web. Act on It."	4 - Imperative, action-oriented	4 - 6 words	5 - Action verb differentiates	4 - Implies analysis + action	4.3
"Your Web Intelligence Platform"	3 - Descriptive but generic	3 - 5 words	2 - "Platform" is overused	5 - Explicit platform claim	3.3
"Intelligence from the Web"	3 - Clean but passive	4 - 4 words	3 - Close to original	5 - Clear scope	3.8
"From URLs to Intelligence"	4 - Clear transformation story	4 - 4 words	4 - Unique framing	3 - Over-emphasizes extraction	3.8
"Every Signal. One Platform."	4 - Rhythmic, balanced	4 - 4 words	3 - Similar to competitor phrasing	5 - "Signal" is broad	4.0

4.2 Complete Homepage Copy

4.2.1 Hero Section Copy

Headline: "Web Intelligence, Delivered."

Subheadline: "Extract data from any URL. Monitor social conversations. Enrich your CRM. Datiq transforms the open web into structured, actionable intelligence for your business. No code required."

Primary CTA: "Start Free" (button, links to signup)

Secondary CTA: "See How It Works" (ghost button, links to demo video or interactive tour)

Social Proof Bar (below hero): "Trusted by 5,000+ users from companies like [logo strip of 4-6 well-known brands]. 10M+ extractions processed. 4.8/5 average rating."

4.2.2 Value Propositions Section

The value propositions section uses a three-column card layout below the hero. Each card represents one of Datiq's core capabilities, with an icon, headline, and two-line description. The copy is designed to communicate the breadth of the platform while maintaining the simplicity that attracted users to the original URL extraction tool.

Table 4.2: Value Proposition Cards

Card	Icon Concept	Headline	Description
Extract	Magnifying glass over a webpage wireframe	Turn Any URL into Data	Paste a URL, get structured data in seconds. Headings, links, contacts, pricing, and AI summaries. No code, no setup.
Social	Pulse/radio wave icon with chat bubbles	Listen to the Web	Track brand mentions across Twitter, Reddit, and news. Real-time sentiment analysis. Know what people say, before they tell you.
Connect	Puzzle piece or hub-and-spoke icon	Integrate Everything	Push extracted data and social insights into HubSpot, Apollo, and your CRM. Bi-directional sync keeps your stack updated.

4.2.3 Social Proof and Trust Section

Below the value propositions, a full-width trust section combines quantitative social proof with qualitative testimonials. The left side shows three large metrics in a horizontal row: "10M+ Extractions Processed," "5,000+ Active Users," and "99.9% Uptime SLA." Each metric uses the platform's accent color for the number and muted text for the label, creating strong visual hierarchy. The right side shows two testimonial cards with the user's name, title, company, and a two-sentence quote. Testimonials should feature specific outcomes rather than generic praise (e.g., "Datiq cut our competitor research time from 4 hours to 20 minutes" rather than "Great tool, highly recommended"). If real testimonials are not yet available, this section can launch with placeholder content from beta users and be updated as reviews accumulate.

4.2.4 Module Overview Section

The module overview section uses an alternating left-right layout to showcase each platform module in detail. Each module gets a dedicated block with a screenshot or illustration on one side and descriptive text on the other. The four modules featured at launch are: Datiq Extract (the original URL extraction tool, repositioned as "Data Extraction" within the broader platform), Datiq Social (the new social listening module), Datiq Connect (the CRM connector hub), and Datiq API (the programmatic access layer). Each block includes a module name, a one-sentence tagline, three bullet-point capabilities, and a "Learn More" link that routes to the module's dedicated landing page or documentation. The alternating layout prevents visual monotony and allows each module to have visual prominence without competing for attention.

4.2.5 Pricing Section and Final CTA

The pricing section displays the three tiers (Free, Starter, Pro) in a horizontal card layout with the Starter tier visually elevated (slightly larger card, accent border, "Most Popular" badge). Each card shows the plan name, monthly price, a list of 5-6 key features with checkmarks, and a CTA button ("Get Started" for Free, "Start Free Trial" for Starter, "Start Free Trial" for Pro). The Enterprise tier is represented by a single call-to-action card: "Need custom limits, SSO, or dedicated support? Talk to our team" with a "Contact Sales" button. Below the pricing, a final CTA section uses a dark background to create visual contrast: "Ready to turn the web into your competitive advantage? Start for free, no credit card required." with a large "Start Free" button. This is the third and final CTA on the page, following the hero CTA and the pricing CTA, ensuring that users who scroll the entire page have multiple conversion opportunities.

4.3 Three Visual Hero Concepts

Each hero concept is defined by its visual composition, layout structure, color treatment, and the psychological impression it creates. All three concepts use the new Datiq brand palette (from Sprint 0.1) and the approved tagline. The concepts are designed to be A/B testable: the team should implement all three and run traffic experiments to identify the highest-converting variant within the first 30 days of launch.

4.3.1 Concept A: "The Dashboard Reveal" (Recommended)

This concept shows a stylized, partially-transparent Datiq dashboard floating above the hero text, creating the impression of the product itself as the hero visual. The dashboard mockup shows three interconnected panels: a URL extraction result on the left (headings, links, AI summary), a social mention feed in the center (sentiment-colored cards), and a CRM sync status panel on the right (HubSpot/Apollo logos with green checkmarks). The dashboard is rendered at a slight 3D perspective angle (rotateX: 5deg, rotateY: -3deg) with a soft drop shadow, giving it depth and tangibility. The background is a very light warm gray (#fafaf8) with a subtle dot grid pattern at 2% opacity, creating a technical but approachable atmosphere. The tagline "Web Intelligence, Delivered." is positioned in the upper-left quadrant (following the F-pattern reading flow) in 48px bold Inter font. The subheadline sits below in 18px regular weight. Both CTAs are positioned in the lower-left, aligned with the tagline. The dashboard visual occupies the right 55% of the viewport. This concept works because it immediately communicates what Datiq does without requiring the user to imagine or infer; the product IS the hero image. It is particularly effective for the "sign up" conversion goal because users can visualize themselves using the product before they register.

4.3.2 Concept B: "The Data Constellation"

This concept uses an abstract visual of interconnected data nodes and lines, representing the web intelligence that Datiq synthesizes from disparate sources. The visual shows 15-20 nodes of varying sizes connected by thin lines, with each node labeled with a data type ("URL," "Tweet," "Review," "CRM Contact," "Price," "Sentiment"). The nodes pulse gently with a subtle CSS animation (scale oscillation between 0.95 and 1.05 over 3 seconds) to convey real-time data flow. The color palette uses the accent color for the largest nodes (representing Datiq's core extraction capability) and muted tones for peripheral nodes (representing the diverse data sources). The background is pure white with the constellation centered in the viewport. The tagline and subheadline are positioned in the center-left with the constellation fading in opacity from left (full opacity near text) to right (30% opacity at edges). This concept is more abstract and brand-forward than Concept A; it works better for the "understand what Datiq is" awareness goal and is well-suited for retargeting campaigns where users already have some familiarity with the product. The risk is that the abstract visual may not communicate the specific functionality as clearly as a product screenshot.

4.3.3 Concept C: "The Before/After Split"

This concept uses a dramatic split-screen composition that contrasts the "before Datiq" state (chaotic, unstructured web data) with the "after Datiq" state (clean, structured intelligence). The left half of the hero shows a collage of raw web elements: a browser screenshot with highlighted text, a Twitter feed with overlapping bubbles, a spreadsheet with messy data, and CRM contact fields with missing information. The right half shows the Datiq-organized version: clean data tables, sentiment charts, enriched CRM cards, and an extraction summary. A bold vertical dividing line separates the two halves, with the Datiq logo positioned at the center of this line, symbolizing the platform as the transformative bridge. The tagline spans the full width above the split: "Web Intelligence, Delivered." The subheadline and CTAs sit below the split. The left half uses a desaturated, slightly blurred treatment (opacity 60%, blur 1px) while the right half is crisp and colorful, creating an immediate visual contrast that communicates value. This concept is the most story-driven and works exceptionally well for first-time visitors who need to understand why Datiq exists and what problem it solves. The risk is that the split composition may feel cluttered on mobile viewports; the responsive version should stack the halves vertically with the "after" state on top.

Table 4.3: Hero Concept Comparison

Dimension	Concept A: Dashboard Reveal	Concept B: Data Constellation	Concept C: Before/After Split
Primary Goal	Conversion (sign-ups)	Awareness (brand understanding)	Education (problem-solution fit)
Visual Style	Product-led, concrete	Abstract, brand-forward	Narrative, contrast-driven
Mobile Adaptation	Dashboard scales down, stacks below text	Constellation reduces to 8-10 nodes	Splits stack vertically, after-state on top
Best Audience	Users familiar with SaaS tools	Return visitors, retargeting	First-time visitors, cold traffic
Development Effort	Medium (Figma dashboard mockup)	Low (CSS/SVG animation)	High (dual illustration, responsive split)
Recommended A/B Test Priority	1st (highest conversion potential)	3rd (brand experiment)	2nd (education for cold traffic)

5. Connector Implementation Plan: HubSpot, Apollo, Stripe, and Calendly

This chapter provides detailed implementation specifications for four priority connectors that extend Datiq's platform utility into the broader business tools ecosystem. Connectors are the force-multiplier that transforms Datiq from a standalone tool into an integrated layer in the user's existing workflow. The four connectors were selected based on three criteria: overlap with Datiq's target customer profile (SMB sales and marketing teams), API maturity and documentation quality, and the frequency of co-usage with web extraction and social intelligence workflows. HubSpot is the most popular SMB CRM with 228,000+ customers; Apollo.io is the leading sales intelligence platform with 275M+ contacts; Stripe is the de facto payment infrastructure for SaaS billing; and Calendly represents the scheduling and meeting-prep workflow where extracted intelligence can automate pre-call research.

5.1 Connector Architecture Overview

All four connectors share a common architectural pattern that ensures consistency, maintainability, and extensibility. The architecture defines a standardized connector interface with five lifecycle methods: `authorize` (initiate OAuth or API key exchange), `test_connection` (verify credentials and API access), `sync` (execute data transfer in one or both directions), `map_fields` (transform between Datiq's data schema and the external system's schema), and `disconnect` (revoke credentials and clean up). Each connector is implemented as an isolated Node.js module with its own configuration, error handling, retry logic, and rate-limit management. A central connector registry manages the list of available connectors, their status, and their per-user authorization state. This abstraction means adding a new connector (e.g., Salesforce, Pipedrive, or

Notion) requires implementing the same five-method interface without modifying the core platform code.

Table 5.1: Connector Architecture Components

Component	Responsibility	Technology	Notes
Connector Registry	Maintains list of available connectors, per-user auth state, sync history	Supabase table (connectors, connector_auth, sync_logs)	Centralized authority for all connector operations
Auth Manager	Handles OAuth 2.0 flows and API key storage with encryption at rest	Supabase Vault (encrypted storage), standard OAuth 2.0 client	Supports both OAuth (HubSpot, Calendly) and API key (Apollo, Stripe)
Sync Engine	Executes data transfer, handles conflict resolution, manages retry logic	Node.js worker queue (BullMQ on Redis)	Bi-directional sync with configurable conflict resolution strategy
Field Mapper	Transforms data between Datiq schema and external system schemas	JSON-based mapping configuration with user-customizable overrides	Pre-built default maps for each connector; users can customize
Rate Limiter	Enforces per-connector API rate limits with token bucket algorithm	Redis-based token bucket, per-connector-per-user isolation	Prevents API throttling and credential suspension
Webhook Receiver	Receives event callbacks from external systems (new contact, deal stage change)	Express.js endpoint with HMAC signature verification	Enables real-time reactive sync rather than polling-only

5.2 HubSpot Connector (Priority: Critical)

HubSpot is the highest-priority connector because CRM integration is the single most requested feature in Datiq's user feedback and the most natural workflow extension: users extract contact and company data from websites, then need to push that data into their CRM for follow-up. The HubSpot connector enables bi-directional synchronization of contacts, companies, and deals between Datiq and HubSpot. In the outbound direction (Datiq to HubSpot), extracted contacts (name, email, phone, title, company) are automatically created or updated as HubSpot contacts with associated company records. In the inbound direction (HubSpot to Datiq), HubSpot contact and deal data enriches Datiq's extraction results, allowing users to see CRM context alongside extracted intelligence. The connector uses HubSpot's OAuth 2.0 authorization flow, which requires Datiq to be listed in the HubSpot App Marketplace (a 1-2 week approval process).

Table 5.2: HubSpot Connector Data Sync Specification

Data Object	Direction	HubSpot API Endpoint	Datiq Field Mapping	Sync Frequency
Contact	Datiq to HubSpot	POST /crm/v3/objects/contacts	email, firstname, lastname, phone, jobtitle, company	On extraction complete (event-driven)
Contact	HubSpot to Datiq	GET /crm/v3/objects/contacts	hubspot_contact_id, lifecycle_stage, deal_pipeline, last_activity	Every 6 hours (scheduled sync)
Company	Datiq to HubSpot	POST /crm/v3/objects/companies	name, domain, industry, description, website, employee_count	On extraction complete (event-driven)
Company	HubSpot to Datiq	GET /crm/v3/objects/companies	hubspot_company_id, type, annual_revenue, owner_email	Every 6 hours (scheduled sync)
Deal	Datiq to HubSpot	POST /crm/v3/objects/deals	dealname (auto-generated), amount, pipeline, stage, close_date	Manual trigger (user clicks "Create Deal")
Note	Datiq to HubSpot	POST /crm/v3/objects/notes	hs_note_body (extraction summary + social intel)	On extraction complete with social context
Social Mention	Datiq to HubSpot	POST /crm/v3/objects/notes (engagement)	Formatted mention with sentiment, source link, and keyword context	Daily digest (batched)

The HubSpot connector's most valuable feature is the automatic enrichment of HubSpot contacts with Datiq extraction data. When a user extracts data from a prospect's website, the connector checks whether a matching contact already exists in HubSpot (by email domain or company domain). If a match is found, the contact record is updated with the extracted data; if no match is found, a new contact and company record are created. This eliminates the manual data entry bottleneck that plagues sales teams and is the primary value proposition that justifies the Starter tier price point. The connector also pushes social intelligence as HubSpot engagement notes, giving sales reps real-time context about prospects' social activity directly in their CRM workflow.

5.3 Apollo.io Connector (Priority: High)

The Apollo.io connector serves a complementary purpose to the HubSpot connector: while HubSpot receives extracted data from Datiq, Apollo provides data to Datiq. Apollo's database of 275 million contacts with verified emails, direct dials, and firmographic data is a powerful enrichment source that transforms Datiq's extraction results from basic web data into comprehensive prospect profiles. The integration uses Apollo's REST API with API key authentication (no OAuth required), which simplifies the implementation compared to HubSpot. The connector supports three primary workflows: contact enrichment (lookup extracted email/domain in Apollo to get full profile), company enrichment (lookup extracted company to get firmographic data like revenue, employee count, tech stack, and industry classification), and sequence sync

(push enriched contacts into Apollo sequences for automated outreach).

Table 5.3: Apollo Connector Workflow Specification

Workflow	Apollo API Endpoint	Input from Datiq	Output to Datiq	Use Case
Contact Enrichment	POST /api/v1/people/match	email or LinkedIn URL from extraction	Full name, title, company, verified email, phone, LinkedIn profile	Enrich extracted contact with complete profile
Company Enrichment	POST /api/v1/organizations/search	company name or domain from extraction	Revenue, employee count, industry, tech stack, funding, description	Add firmographic context to extracted company data
Sequence Push	POST /api/v1/sequences/add_contacts	Enriched contact IDs from Datiq	Contact added to Apollo outreach sequence	Automated follow-up after extraction
Email Verification	GET /api/v1/people/match?verify_email=true	Email address from extraction	Verification status (valid, invalid, catch-all), deliverability score	Filter out invalid emails before CRM push

The Apollo connector creates a powerful multi-step workflow: a user extracts a prospect list from a conference website using Datiq Extract, the Apollo connector enriches each contact with verified emails and firmographic data, and the HubSpot connector pushes the enriched profiles into the CRM as contacts and companies. This three-step pipeline (Extract, Enrich, Sync) is the killer use case that differentiates Datiq from both standalone extraction tools (which lack enrichment) and standalone sales intelligence tools (which lack extraction). The combined workflow saves sales teams an estimated 3-4 hours per prospect list compared to manual research and data entry. Apollo's API pricing is based on the user's Apollo subscription tier; Datiq passes through the user's own Apollo API key, ensuring that Datiq does not incur per-call costs and the user's existing Apollo credits are used.

5.4 Stripe Connector (Priority: Critical for Billing)

The Stripe connector serves a dual purpose: it powers Datiq's own subscription billing (as described in Chapter 3) and, in Phase 2, will enable users to extract pricing data from competitor Stripe payment pages and sync it with their own pricing analysis workflows. In Phase 1, the Stripe integration is internal-facing only: it manages subscription creation, plan changes, usage metering, and invoice generation. The implementation uses Stripe's Checkout Sessions for one-time plan purchases and the Customer Portal for self-service plan management (upgrade, downgrade, cancel, update payment method). Usage-based billing leverages Stripe's Metered Billing API, which accumulates usage records throughout the billing cycle and generates invoices based on per-unit pricing.

Table 5.4: Stripe Connector Implementation Specification

Stripe Feature	API / Method	Datiq Usage	Implementation Notes
Checkout Session	stripe.checkout.sessions.create()	New subscription creation (Free to Starter/Pro)	Redirect-based flow; supports coupon codes and trial periods
Customer Portal	stripe.billingPortal.sessions.create()	Self-service plan management	Pre-configured URL; users manage payment methods and view invoices
Metered Billing	stripe.subscriptionItems.createUsageRecord()	Track extraction credits and social mentions per cycle	Usage records uploaded in real-time; aggregated at billing cycle end
Webhook Events	POST /webhook (customer.subscription.updated, invoice.paid)	Sync subscription status with Datiq user record	HMAC-SHA256 signature verification; idempotency key handling
Tax Calculation	stripe.tax.calculations.create()	Auto-calculate VAT/GST for EU and Indian customers	Requires Stripe Tax registration; uses customer location

5.5 Calendly Connector (Priority: Medium, Phase 2)

The Calendly connector is planned for Phase 2 (Weeks 13-18) but is included in this specification because the integration design has implications for the connector architecture built in Phase 1. The Calendly connector's value proposition is meeting intelligence automation: when a prospect books a meeting through Calendly, the connector automatically triggers a Datiq extraction of the prospect's company website, enriches the data via Apollo, and pushes a pre-meeting briefing card to the user's email or Slack. This transforms the 15-20 minutes of manual pre-call research that sales reps typically do into an automated, instantaneous workflow. The connector uses Calendly's Webhook API (v2) to receive booking events and the Users API to read meeting details (attendee name, email, meeting type, questions and answers).

Table 5.5: Calendly Connector Workflow Specification

Trigger	Calendly Event	Datiq Action	Output
Meeting Booked	invitee.created webhook	Extract attendee company website, enrich via Apollo, generate briefing	Pre-meeting briefing email with company overview, key contacts, recent news, social sentiment
Meeting Rescheduled	invitee.rescheduled webhook	Update briefing with any new data (if >24h since original)	Updated briefing email (only if significant new data found)
Meeting Canceled	invitee.canceled webhook	Log cancellation, archive briefing, no further actions	Internal notification only
Daily Digest	Scheduled (Cron, 8 AM user timezone)	Aggregate all upcoming meetings for the day, generate batch briefings	Single digest email with all day's meeting briefings

The Calendly connector's pre-meeting briefing feature is the most compelling user-facing workflow in the entire connector ecosystem. Imagine a sales rep who has 6 meetings in a day: without Datiq, they spend 1.5-2 hours researching each company. With the Calendly connector, they wake up to a single email containing a structured briefing for each meeting: company overview, key decision-makers, recent news and social sentiment, competitive positioning, and suggested talking points. This feature alone can justify the Pro tier subscription for individual sales professionals and is a powerful differentiator against competitors like Apollo (which offers company data but not automated pre-meeting briefings) and Crayon/Klue (which offer competitive intelligence but not meeting-specific automation). The connector architecture built in Phase 1 (specifically the webhook receiver and field mapper) directly enables this workflow, validating the upfront investment in a generic connector abstraction layer.

5.6 Connector Implementation Timeline and Dependencies

Table 5.6: Connector Implementation Timeline

Connector	Phase	Sprint	Est. Effort	External Dependencies	Revenue Impact
HubSpot	Phase 1	Sprint 1.2 (Wk 7-8)	8 person-days	HubSpot App Marketplace approval (1-2 weeks)	High - primary conversion driver for Starter tier
Apollo	Phase 1	Sprint 1.2 (Wk 7-8)	6 person-days	Apollo API key (user-provided)	High - enrichment workflow justifies Pro tier
Stripe (billing)	Phase 1	Sprint 0.2 (Wk 3-4)	5 person-days	Stripe account, Tax registration (optional)	Critical - enables all revenue collection
Calendly	Phase 2	Sprint 2.1 (Wk 13-14)	5 person-days	Calendly Webhook API access	Medium - retention driver for Pro/Enterprise
Salesforce (planned)	Phase 2	Sprint 2.2 (Wk 15-16)	8 person-days	Salesforce AppExchange listing	High - expands market to enterprise
Pipedrive (planned)	Phase 2	Sprint 2.3 (Wk 17-18)	5 person-days	Pipedrive API access	Medium - popular CRM in EMEA markets

The total connector development effort across Phase 1 and Phase 2 is approximately 37 person-days, with the HubSpot and Salesforce connectors being the most complex due to their OAuth flows, extensive data models, and marketplace approval requirements. The connector architecture investment in Phase 1 (approximately 3-4 additional person-days for the abstraction layer) reduces the per-connector implementation time by roughly 30% for Phase 2 connectors, as the auth manager, sync engine, and field mapper are already built and tested. The recommended sequencing prioritizes billing (Stripe) first because it enables revenue collection, then CRM connectors (HubSpot, Apollo) because they drive conversion, and finally workflow connectors (Calendly) because they drive retention and expansion revenue.