

DatIQ

Prioritised Features Release Roadmap

MVP → v2.1 → v5.0+ · A Model Council Synthesis

Prepared by	4-Member AI Model Council (Growth & Conversion · Market Strategy · Retention & UX · Platform Engineering)
Evidence base	Live review of datiq.app (v2.0) + 6 independent model analyses: Gemini, Perplexity, ChatGPT, Grok, DeepS
Scope	62 candidate features scored on 6 parameters, bundled into 7 releases
Date	July 16, 2026
Companion file	Prioritised Features Release Roadmap – Backlog Tracker.xlsx (editable scoring model)

1. Executive Summary

DatIQ (datiq.app, v2.0) is a zero-code web extraction and AI enrichment platform — paste any public URL and receive headings, links, contacts, pricing tiers, an AI summary, custom plain-English field extraction, domain maps and SEO content briefs in about 30 seconds. Its single-URL 'magic moment' is genuinely differentiated: no competitor bundles this breadth per URL behind so little friction. The market tailwind is strong (web scraping software ~USD 1.0–1.6B in 2026, 13–19% CAGR; AI-driven extraction growing faster).

The council's central diagnosis, shared across all six model reports and all four council members, is that DatIQ creates value and then lets it evaporate. The product is stateless (no visible history or projects), its results are trapped on-screen (no export, no shareable URLs), and the client-rendered single-page app is effectively invisible to search engines — no indexed subpages, pricing, docs or changelog. These three leaks — value, memory and visibility — cap conversion, retention and acquisition simultaneously, and fixing them costs far less than any headline feature.

Accordingly, the roadmap front-loads a small, cheap 'Convert & Keep' release (export, accounts + history, transparent pricing, extract-similar, paywall flow — ~5.5 person-months), follows with an acquisition flywheel (shareable SSR permalinks, programmatic SEO, persona onboarding, referral loop), then builds the habit engine (collections, batch depth, scheduled monitoring with change alerts — the only feature endorsed unanimously by every source and every council member), the developer/agent ecosystem (API, SDK, MCP server, Zapier, template library), an intelligence & quality tier (self-healing extraction, confidence scoring, verified emails, tech-stack detection, domain dashboards), teams & enterprise, and finally the agentic vision (DatIQ Agent, chat-over-library, lookalike discovery) — deliberately parked until reliability infrastructure and unit economics support it.

Two of the eight supplied reports analysed the wrong product (Copilot reviewed a dating app; Meta AI reviewed a spreadsheet-analysis app on datiq.netlify.app). The council screened rather than discarded them, salvaging transferable ideas (onboarding tooltips, deep-linked re-engagement, prompt chips, privacy reassurance, localization) while excluding everything product-specific. No differentiating idea from any source was dropped without a recorded verdict: all 62 candidates appear in the scored backlog.

2. Product Baseline — What Is Live Today (verified 16 July 2026)

Live modules: Heading Structure, Every Link, AI Summary, Custom Extraction (plain English), Domain Mapping, Contacts & Emails, Content Generation (SEO outlines/briefs), Pricing Extraction. Also live: Batch and Schedule tabs, Advanced options, quick-example chips, trial mode (10 extractions · 5 batch runs), Pro \$29/mo (250 extractions), Agency \$199/mo (unlimited). Site metrics shown: 11+ teams & researchers, 138+ extractions, 30s average time to insight, 7 export & enrichment types.

Because Batch and Schedule already exist in early form, the roadmap treats them as features to deepen (CSV upload, results tables, diff-based change alerts) rather than to build from scratch — a correction to several source reports that assumed they were absent.

Structural weaknesses confirmed: client-rendered SPA invisible to crawlers; no shareable result permalinks; no visible export; no account memory; no API, extension or integrations; no team features; opaque pricing (uncrawlable); tiny usage base (138 extractions) meaning acquisition and activation dominate every other concern.

3. Evidence Base and Source Vetting

Source	Focus / signature recommendations	Council verdict
Gemini (2 files)	Schema auto-detection, SDK code exporter, MCP server (top 3); change monitoring; interactive action sandbox; autonomous agentic crawler; competitive matrix vs Firecrawl/ScrapeOps/Browse.ai/Browserbase.	On-product. MCP/SDK adopted (R4); sandbox legal-gated (R7); agent deferred (R7).
Perplexity-style RICE++ report (md + pdf)	27 features RICE++-ranked; homepage demo (43.2), export (41.0), saved history (38.9) on top; full sprint roadmap; PLG loops; persona landing pages.	On-product; broadest single source. Heavily adopted across R1–R5.

Source	Focus / signature recommendations	Council verdict
ChatGPT-style weighted report	SSR + deep-linkable permalinks as #1 (4.15); programmatic SEO; changelog + Product Hunt; staged benchmarks (activation 20–40%, free→paid 3–8%). Flagged SEO invisibility as the single biggest gap.	On-product. Adopted; its metrics became release KPIs.
Grok (shared link)	Browser extension (4.5), persistent library + auto-tagging (4.3), dashboard/history feed (4.2), AI chat over saved library, page snapshots/anti-link-rot.	On-product. Library/watchlist ideas shaped R3; chat-over-library in R7.
DeepSeek (shared link)	'Extract Similar' CTA (9.5), role templates (9.0), export, monitoring (9.0), extension (9.0); pricing intel (trial limits, \$29/\$199 tiers); Teams-tier + metered-API monetisation.	On-product; only source with in-app pricing detail. Monetisation section adopted (F42, FA3).
Meta AI (shared link)	Analysed datiq.netlify.app — a spreadsheet-insights product. Smart input router, example cards + prompt chips, shareable insight pages, gallery/remix, India-first localization, OCR.	WRONG PRODUCT — screened. Transferable ideas salvaged (F05, F14, F44, gallery pattern); spreadsheet-specific items excluded.
Copilot (pdf)	Analysed a dating app. Onboarding tooltips, deep-linked notifications, QR sharing, analytics dashboards, A/B discipline, safety/verification.	WRONG PRODUCT — screened. Salvaged F15, F49, F50 patterns; all dating mechanics excluded.
Live site + council research	Verified v2.0 module list, Batch/Schedule presence, trial limits, stats; competitor pricing (Firecrawl, Browse AI, Clay, Hunter, Kadoa, Thunderbit).	Primary ground truth; overrode report assumptions where they conflicted.

4. Methodology — Council Composition and Scoring Model

Four council members deliberated independently from the same consolidated briefing, then their positions, objections and cross-challenges were synthesised by a chair:

Member A — Growth & Conversion. Free-to-paid funnels, activation, PLG loops, pricing. Championed export, accounts, pricing transparency, payroll flow, referral and public-report mechanics.

Member B — Market & Competitive Strategy. Positioning vs Firecrawl/Browse AI/Clay/Hunter/Kadoa, segment economics, GTM. Championed permalinks, programmatic SEO, MCP, template library, comparison pages, lookalike discovery.

Member C — Engagement, Retention & UX. Habit loops, D1/D7/D30, discoverability, perceived quality. Championed history, collections, watchlist home, monitoring, confidence badges, digest emails.

Member D — Platform Engineering & Quality. Feasibility, dependencies, reliability, cost, compliance. Championed anti-bot/JS infra, job queue + cache, self-healing, compliance engine; vetoed premature agentics and credentialed scraping.

Scoring model

Each feature scored 0–5 on six parameters mapped to the stated business goals: CONV conversion-rate impact (weight 22%), REACH market reach/acquisition (16%), ENG engagement & retention (20%), EQ efficiency & output quality (12%), DIFF differentiation (15%), MON monetisation fit (15%). Effort estimated in person-months.

Value = $\Sigma(\text{score} \times \text{weight})$ · **Priority Index** = $\text{Value} \times (1 + 0.05 \times \text{Council Votes}) \div \sqrt{\text{Effort}}$

Council Votes (0–4) counts members placing the feature in their top-12. The square root of effort prevents large strategic bets from being annihilated by cheap cosmetics while still rewarding quick wins. Scores are member averages where scored, chair-assigned from source evidence otherwise. The companion Excel recalculates everything if weights or scores are edited. Rankings guide sequencing, but release assignment also honours hard dependencies (accounts → collections → teams; infra → monitoring GA; API → SDK/MCP).

5. Council Debate — Key Tensions and Resolutions

Debate	Positions	Resolution
Ship viral permalinks (F04) first vs. fix retention first	Growth (A) + Market (B) wanted F04 in first release; Retention (C) objected — sharing into a stateless product amplifies churn; Engineering (D) flagged DMCA + cache dependency.	RESOLVED: F04 ships in R2, immediately after accounts/history (R1) and with the FD2 result-cache and a takedown process. Viral loop lands on a product that can retain the users it acquires.
Monitoring now vs. infrastructure first	All four members back monitoring (F17) — the only unanimous feature. Engineering (D) insisted a monitor that silently breaks on redesigns is a churn event with a paying customer attached.	RESOLVED: F17 beta in R3 on a curated domain set, GA in R4 gated on anti-bot/JS infra (F36) and retry ladder (FD1 partial). C accepted the sequencing.

Debate	Positions	Resolution
Developer lane (API/MCP) vs. non-developer focus	Growth (A) challenged the dev play — Firecrawl owns that lane at \$16/mo. Market (B) countered that MCP is distribution, not a head-on dev fight; Engineering (D) showed F23 is a thin wrapper once F21 exists.	RESOLVED: API+MCP in R4 as a distribution channel with persona-flavoured tools (extract_pricing, find_contacts) rather than a raw-scraping API price war. Primary positioning stays no-code.
Autonomous agent as flagship vs. science project	Gemini + Perplexity reports pitched agentic crawling as the future; every council member voted DELAY/KILL for 12 months — cost 20-50x per run, compounding error rates, no reliability layer yet.	RESOLVED: F34 parked in R7 (2028+), explicitly gated on F35 self-healing, F36 infra, and positive unit economics. MCP (F23) carries the 'agent-ready' narrative meanwhile.
Team features now vs. teams pricing first	DeepSeek + Grok pushed team workspaces early; Growth (A): with 11 teams total there is no multiplayer demand yet — 'sell seats, then build for them.'	RESOLVED: Teams *pricing tier* (F42) in R4; full workspaces/RBAC (F31) in R6 after collections (F03) prove shared-work demand. Dependency chain F02→F03→F31 respected.
Auth/paywalled-page extraction (F33)	Gemini ranked the action sandbox high (10.0); Engineering (D) flagged CFAA/ToS/GDPR exposure and credential-vault security burden a small team cannot carry.	RESOLVED: KILLED from the 18-month roadmap; revisit post-SOC2 (F41) with legal counsel only.
Wrong-product source reports	Copilot analysed a dating app; Meta AI analysed a spreadsheet-analysis product (datiq.netlify.app). Council screened both for transferable ideas rather than discarding.	RESOLVED: Salvaged — onboarding tooltips (F15), re-engagement deep links (F49), analytics discipline (F50) from Copilot; prompt chips (F05), privacy reassurance (F14), localization (F44), micro-tool/gallery patterns from Meta AI. All product-specific recommendations (swipe mechanics, spreadsheet connectors) excluded.

6. Top 20 Features by Priority Index (all releases)

#	ID	Feature	Release	Value	Priority	Effort (pm)
1	F01	One-click Export (CSV / JSON / Clipboard / Google Sheets)	R1 · v2.1	3.18	5.77	0.4
2	F07	"Extract Similar" Post-result CTA	R1 · v2.1	3.01	5.24	0.4
3	F13	Transparent Pricing Page + Tier-Feature Mapping	R1 · v2.1	2.03	4.47	0.25
4	F22	SDK Code-snippet Exporter (Python / Node / cURL from any manual run)	R4 · v3.5	2.92	4.34	0.5
5	FA1	Free-tier "Public Report" Quota Mechanic	R2 · v2.2	3.45	4.18	0.75
6	F24	LLM-ready Markdown Export (RAG/agent pipelines)	R4 · v3.5	2.71	4.02	0.5
7	F23	MCP Server (DatIQ as native tool for Claude/GPT/Cursor agents)	R4 · v3.5	3.51	3.45	1.25
8	F14	Privacy & Trust Messaging + Data-retention Statement	R1 · v2.1	1.64	3.44	0.25
9	FA2	Referral Credits Loop (give 25 / get 25)	R2 · v2.2	2.30	3.42	0.5
10	F02	Accounts + Saved Extraction History	R1 · v2.1	3.37	3.40	1.3
11	F06	Persona Onboarding + Recipe Packs (Sales Pack, CI Pack, SEO Pack...)	R2 · v2.2	3.35	3.30	1.25
12	FB1	Competitor Comparison Pages + Battle-card Generator (vs/firecrawl...)	R2 · v2.2	3.07	3.22	1
13	F04	Shareable Result Permalinks + SSR ("Powered by DatIQ" loop)	R2 · v2.2	3.56	3.19	1.5
14	F19	Zapier / Make / Webhooks	R4 · v3.5	2.79	3.07	1
15	F17	Scheduled Monitoring + Change Detection Alerts (email/Slack)	R3 · v3.0 (beta) → R4 GA	3.79	2.94	2.4
16	F09	AI Output Feedback Widget (thumbs up/down + comment)	R1 · v2.1	1.94	2.88	0.5
17	FC3	Weekly "Your Data Changed" Digest (auto-diff summary email)	R4 · v3.5	2.67	2.80	1
18	FA3	In-product Paywall / Quota-exhaustion Upgrade Flow + Annual Plans	R1 · v2.1	1.85	2.75	0.5
19	F25	Prebuilt Schema / Template Library (e-com, jobs, news, SaaS pricing...)	R4 · v3.5	3.66	2.72	2
20	F08	Preset Schema Auto-detection from URL Patterns	R4 · v3.5	2.65	2.65	1

Highest-Value features irrespective of effort: F17 Scheduled Monitoring + Change Alerts (3.79 — the only unanimous council pick), F25 Template Library (3.66), F04 Shareable Permalinks (3.56), F23 MCP Server (3.51), FA1 Public-Report Quota (3.45). These anchor releases R2–R4.

7. Release Roadmap — Full Prioritised Backlog by Release

All 62 features are assigned; nothing from any source was dropped without a verdict. Within each release, features are listed by Priority Index. Effort totals assume dev+design+QA person-months; R1+R2 fit a 2–3 engineer team in one quarter.

R1 · v2.1 — “Convert & Keep” · 0–6 weeks (Aug 2026) · 10 features · ~5.5 person-months

Objective: Stop the leaks: make value portable, persistent and purchasable. Fix activation→conversion before buying traffic.

Exit criteria / KPIs: D1 activation ≥40% (2+ extractions + 1 export); free→paid ≥3%; baseline telemetry live

ID	Feature & council note	Pri	Effort	Sources
F01	One-click Export (CSV / JSON / Clipboard / Google Sheets) Table stakes across every competitor; #1 trial-killer today — value dies on-screen. Unanimous council support.	5.77	0.4	Perplexity, ChatGPT, Grok, DeepSeek, Gemini
F07	"Extract Similar" Post-result CTA Cheapest session-depth multiplier; turns 1 extraction into 4, burns trial quota faster, hits paywall sooner.	5.24	0.4	DeepSeek (9.5/10 top pick); Council A/C

ID	Feature & council note	Pri	Effort	Sources
F13	Transparent Pricing Page + Tier-Feature Mapping Buyers comparison-shop ruthlessly; an uncrawable pricing page loses deals before they start.	4.47	0.25	ChatGPT, DeepSeek; Council A/B
F14	Privacy & Trust Messaging + Data-retention Statement 'Data deleted after Xd, no training on your data' under the input box; table stakes before Agency/API deals.	3.44	0.25	Meta AI (transferable), Council D
F02	Accounts + Saved Extraction History Substrate for every later feature (collections, monitoring, API keys, teams). Free tier capped at 10 items = upgrade trigger.	3.40	1.3	All 5 on-topic reports; Council A/C/D
F09	AI Output Feedback Widget (thumbs up/down + comment) Instruments perceived output quality — the silent retention variable; feeds extraction-quality improvement loop.	2.88	0.5	Perplexity; Council C
FA3	In-product Paywall / Quota-exhaustion Upgrade Flow + Annual Plans Upgrade screen at the exact moment of trial exhaustion, showing which paid feature completes the current task; annual anchoring.	2.75	0.5	Council A (new)
F15	Onboarding Tooltips / First-run Product Tour Clarifies the 8 extraction modes on first visit; pairs with F06 recipe chips.	2.05	0.5	Copilot (transferable), Perplexity
F50	Product Analytics + A/B Instrumentation (foundational) Non-negotiable: every later prioritisation claim is untestable without funnel + retention telemetry.	1.80	0.9	Copilot (transferable), Council C/D
F10	Keyboard Shortcuts / Power-user Mode Low-cost delight for analysts and power users.	1.64	0.5	Perplexity

R2 · v2.2 — “Launch & Be Found” · 6–12 weeks (Sep–Oct 2026) · 8 features · ~7.25 person-months

Objective: Acquisition flywheel: every result a shareable indexed page; persona routes; public launch moment.

Exit criteria / KPIs: Indexed pages >50; organic visits baseline; referral K-factor >0.15; PH launch top-5 of day

ID	Feature & council note	Pri	Effort	Sources
FA1	Free-tier "Public Report" Quota Mechanic Unlimited free extractions if published as public indexed DatIQ-branded pages; private runs consume quota. Feeds F04/F11 content flywheel.	4.18	0.75	Council A (new)
FA2	Referral Credits Loop (give 25 / get 25) Triggered at quota-exhaustion moment; classic PLG, near-zero build.	3.42	0.5	Council A (new)
F06	Persona Onboarding + Recipe Packs (Sales Pack, CI Pack, SEO Pack...) Routes SDRs to Contacts and analysts to Pricing in <60s; blank-box tools bleed activation.	3.30	1.25	Perplexity, DeepSeek (9.0), Browse AI benchmark; Council A/C
FB1	Competitor Comparison Pages + Battle-card Generator (/vs/firecrawl...) Comparison keywords = highest-intent SaaS traffic; 'compare 2 competitor URLs → diff table' demos the product while ranking for it.	3.22	1	Council B (new)
F04	Shareable Result Permalinks + SSR ("Powered by DatIQ" loop) Every extraction becomes an indexed, shareable landing page — fixes SEO invisibility + creates the viral loop. Requires FD2 cache; ship after accounts (Council C condition) with DMCA takedown process (Council D condition).	3.19	1.5	ChatGPT (#1, 4.15), Perplexity, Meta AI; Council A/B
F11	Programmatic SEO + Persona Landing Routes (/for-sales, /extract-pricing...) Near-zero SEO footprint in a 13-19% CAGR category; long-tail intent pages the single home URL can never capture.	2.62	1	ChatGPT (4.00), Perplexity; Council A/B
F05	Homepage Animated Demo + Example Gallery + Prompt Chips Pre-filled animated extraction on a known URL = instant aha with zero signup friction; prompt chips teach what to ask.	2.40	0.75	Perplexity (#1, 43.2), Meta AI, DeepSeek
F12	Public Changelog + Product Hunt / Social Launch v2 needs a launch moment; changelog builds trust and re-engagement.	2.12	0.5	ChatGPT (3.75); Council A

R3 · v3.0 — “Habit & Reliability” · Q4 2026 · 9 features · ~15.65 person-months

Objective: Turn one-shot use into weekly habit; harden the engine (JS rendering, queue, compliance) that everything else depends on.

Exit criteria / KPIs: D7 retention 2× baseline; extractions/session ≥3; monitoring beta on 500 URLs; fetch success ≥95%

ID	Feature & council note	Pri	Effort	Sources
F17	Scheduled Monitoring + Change Detection Alerts (email/Slack) Converts one-shot tool into recurring-value subscription; the D30 retention engine and strongest \$29/mo justification. Schedule tab exists — add diffing + alerting. GA gated on F36.	2.94	2.4	ALL 5 reports + all 4 council members (only unanimous feature)
FC1	Watchlist Home (logged-in dashboard: monitored URLs, deltas, re-run) The 'Bloomberg-terminal opening screen' for web data; gives history + monitoring a daily-open surface.	2.64	1.5	Council C (new); Grok (dashboard/history feed)
F16	Batch Enhancements (CSV URL upload, results table, parallel runs) Forces the queue/worker/retry architecture needed for monitoring + API anyway; closes gap vs Hunter bulk & Clay.	2.54	1.75	All reports; batch exists in trial — deepen it
F18	Google Sheets / Airtable / Notion Export Integrations Where extracted data actually lives for non-technical personas; cheapest 'integrations' checkbox.	2.33	1.5	ChatGPT, Grok, DeepSeek
F36	Anti-bot / Proxy / JS-rendering Infrastructure ~40% of target pages are JS-heavy or bot-protected; existential dependency for batch-at-scale, monitoring GA and API credibility.	2.08	2.5	Gemini, DeepSeek, Bright Data benchmark; Council D (#1)
FD3	robots.txt / ToS Compliance Engine + Per-domain Rate Limiting Prevents shared-infra IP bans (one abusive user blackholes everyone); the compliance story enterprise buyers ask about first.	2.06	1	Council D (new)
F49	Deep-linked Re-engagement Emails ('your schedule ran / data changed') Free D7 lift once accounts + monitoring exist; deep links resume exact context.	2.04	1	Copilot (transferable); Council C
F03	Collections / Projects Workspace Named research projects (e.g. 'Q3 Competitor Research'); organised accumulation = switching cost; data model for teams later.	1.90	2.5	Perplexity, Grok, DeepSeek (Intelligence Hub); Council C
FD2	Idempotent Job Queue + Result Cache + Dedup Same URL twice in an hour = one fetch; cuts infra cost 30-50%, makes permalinks instant and monitoring predictable.	1.56	1.5	Council D (new)

R4 · v3.5 — “Ecosystem & Developers” · Q1 2027 · 9 features · ~10.95 person-months

Objective: Open the platform: API, SDK, MCP, Zapier, template library; monetisation architecture (Teams tier, metered API).

Exit criteria / KPIs: API signups; MCP installs; Zapier listing live; WAU/MAU ≥40%; free→paid ≥5%

ID	Feature & council note	Pri	Effort	Sources
F22	SDK Code-snippet Exporter (Python / Node / cURL from any manual run) Pure templating over the API; converts UI users into API users.	4.34	0.5	Gemini (top-3, 16.0); Council D
F24	LLM-ready Markdown Export (RAG/agent pipelines) Near-free (DOM already parsed); rides the RAG wave; captures the AI-engineer persona.	4.02	0.5	Perplexity, ChatGPT, Firecrawl benchmark; Council D
F23	MCP Server (DatIQ as native tool for Claude/GPT/Cursor agents) Thin wrapper over the API; the agent ecosystem is 2026's fastest-growing distribution channel. Ship persona-flavoured tools (extract_pricing, find_contacts).	3.45	1.25	Gemini (15.63), Perplexity, ChatGPT; Council B/D
F19	Zapier / Make / Webhooks Zapier directory listing is itself an acquisition channel; 'URL extracted → push to Sheet/Notion/HubSpot' = live data pipeline.	3.07	1	Perplexity, ChatGPT, Grok, DeepSeek; Council A/B
FC3	Weekly "Your Data Changed" Digest (auto-diff summary email) Persona-valuable habit loop (not gamification fluff); extends monitoring value between sessions.	2.80	1	Council C (new); Perplexity digest idea
F25	Prebuilt Schema / Template Library (e-com, jobs, news, SaaS pricing...) Each template doubles as an indexed SEO landing page (Browse AI's engine); typed JSON output serves developers too.	2.72	2	All reports; Browse AI (150+) benchmark; Council B
F08	Preset Schema Auto-detection from URL Patterns Regex + semantic URL analysis auto-suggests product/job/article schema; reduces configuration friction.	2.65	1	Gemini (top pick, 16.0)

ID	Feature & council note	Pri	Effort	Sources
F21	Public REST API + Key Management + Usage Metering Unlocks dev/AI segment + usage-based revenue; B2B2B distribution. Needs F02 + rate limiting.	2.57	2.2	Perplexity, ChatGPT, DeepSeek, Gemini; Council B/D
F42	Teams Pricing Tier + Usage-based Add-ons + Paid API Tier Teams \$79-99/mo, metered API, enterprise custom; replaces 'Agency unlimited' abuse magnet. Sell seats before building full RBAC.	1.87	1.5	DeepSeek (monetisation section); Council A/B

R5 · v4.0 — “Intelligence & Quality” · Q2–Q3 2027 · 9 features · ~19.7 person-months

Objective: Differentiate on being right: self-healing, confidence scoring, verified emails, tech-stack + domain intelligence; extension distribution.

Exit criteria / KPIs: Extraction accuracy ≥97% on template set; verified-email credits revenue; extension WAU; D30 retention 2×

ID	Feature & council note	Pri	Effort	Sources
F26	Tech Stack Detection (BuiltWith/Wappalyzer-style) High-value intel for SDRs (what CRM do they use?), VCs and devs; unique vs current no-code competitors.	2.57	1.5	Perplexity (7.7)
F27	Domain Intelligence Dashboard (live company one-pager) All URLs + structure + contacts + tech + latest content for any domain — a mini-Diffbot at 1/50th the price; serves VC, SDR and CI at once.	2.52	3	Perplexity, Grok; Council B (category-creating framing)
F20	Chrome Extension (one-click extract on current tab) Chrome Web Store = free discovery surface; ambient retention; Thunderbit/Hunter prove the wedge.	2.49	2.7	All 5 reports (Grok #1, DeepSeek 9.0); Council A/B/C
FD1	Extraction Reliability Layer (status page, confidence scores, retry ladder) fetch → headless → proxy fallback + per-result confidence; the 'it just works' difference between toy and tool.	2.43	2	Council D (new)
F35	Self-healing Extraction + Validation Layer Semantic re-location of fields when layouts change; drift detection; silent-wrong-data kills trust faster than errors.	2.16	3	ChatGPT, Gemini, Kadoa benchmark; Council D
FB2	Free Public Micro-tools (email finder, pricing extractor, sitemap visualiser) Each micro-tool = a distinct indexed page funnelling to signup; Hunter built its entire top-of-funnel this way.	2.14	1.5	Council B (new); Hunter's playbook
FC2	Per-field Confidence Badges + Inline "Fix this field" Corrections One-click correction retrains the schema; builds trust and a proprietary quality dataset.	2.05	2	Council C (new)
F28	Email Verification / Deliverability Scoring A verified email is worth 5-10x a raw scraped address; credit-based monetisation aligned with Hunter's model.	1.84	3	Perplexity, ChatGPT, Hunter benchmark
F48	User Analytics Dashboard (usage, quota, insights unlocked) Transparency into consumption; natural upgrade prompt surface.	1.73	1	Grok, Copilot (transferable)

R6 · v4.5 — “Teams & Enterprise” · Q4 2027 · 5 features · ~15 person-months

Objective: Multiplayer + enterprise readiness: workspaces, CRM, SOC2 (trigger-based), CI feed as vertical product.

Exit criteria / KPIs: Multi-seat accounts ≥15% of revenue; first enterprise deals; NRR >110%

ID	Feature & council note	Pri	Effort	Sources
F30	CRM Integrations (HubSpot first, then Salesforce/Pipedrive) 'URL → enriched HubSpot record' is a Clay-lite pitch at a tenth the price; SDR willingness-to-pay lives here.	2.19	2.5	Perplexity, DeepSeek; Council B
F32	Competitor Change-Intelligence Feed (curated CI dashboard) Vertical SaaS product for PM/strategy teams: pricing updates, feature launches, team changes in one feed; builds on monitoring.	1.56	4	Perplexity (4.5); Browse AI mid-market benchmark
F44	Localization / Multilingual Extraction + India-first GTM Nobody in the category does India-first; large underserved SMB/analyst market; pairs with regional payment methods.	1.56	2.5	Meta AI (transferable: Hinglish, regional pricing)
F31	Team Workspaces + RBAC + Shared History Primary unlock for team/enterprise pricing; deliberately sequenced after collections prove multiplayer demand.	1.44	3	All reports; Council consensus: AFTER F02→F03 and Teams-tier demand proof
F41	SOC 2 / GDPR Compliance Program Trigger-based: pursue when enterprise deals or design-partner pipeline justify the spend.	0.79	3	ChatGPT; Council D

R7 · v5.0+ — “Agentic Vision” · 2028+ · 12 features · ~41.5 person-months

Objective: Category creation: autonomous research agent, chat-over-library, lookalikes, waterfall enrichment, knowledge graph.

Exit criteria / KPIs: Agent unit economics positive; library-chat WAU; ARPU expansion

ID	Feature & council note	Pri	Effort	Sources
FB3	Lookalike Company Discovery ('find 25 similar companies') Turns one extraction into a prospecting list; massively raises SDR willingness-to-pay.	2.42	3	Council B (new); Clay/Ocean.io benchmark
F38	AI Chat over Saved Library (ask questions across all saved extractions) Makes the accumulated library compounding, not just stored; RAG over user's own research corpus.	1.81	4	Grok (signature idea); Council C
F34	"DatIQ Agent" — Autonomous Multi-step NL Research Crawler 'Find CEO, pricing, tech stack for these 30 companies' autonomously. Gated on F35+F36 reliability and unit economics (20-50x per-run cost).	1.49	6	Gemini, Perplexity (Claygent parity); Council verdict: unanimous DELAY 12 months
F46	Milestones / Gamification ('50 companies analysed!') Folded into FC3 digest; badges on a tool with no habit loop is decoration.	1.40	1	Perplexity, Copilot (transferable); Council B: 'retention theatre'
F37	Full-page Snapshots / Archiving (anti-link-rot) Historical archive of monitored pages; useful for CI/legal use cases.	1.37	2	Grok
F29	Waterfall Enrichment (multi-source sequential fill) Start with 2-source fallback (own extraction + email lookup); full waterfall is a Clay-scale project.	1.36	5	Perplexity, DeepSeek; Clay's signature innovation
F47	Split-screen Live Render Preview Visual confirmation of extraction accuracy; nice-to-have after confidence badges (FC2).	1.35	1.5	Gemini
F33	Interactive Action Sandbox (logins, clicks, auth/paywalled pages) CFAA/ToS/GDPR minefield + credential-vault burden; revisit only post-F41 with counsel.	1.29	5	Gemini, Firecrawl/Browserbase benchmark; Council D: KILL until legal cleared
F43	Public Gallery / Remix of Shared Reports Network-effect play once shared permalinks (F04) reach volume.	1.18	2	Meta AI (transferable)
F40	Template Marketplace (creator monetisation) Marketplace economics need thousands of MAU creators; ship first 100 templates in-house (F25) instead.	0.98	4	DeepSeek, Meta AI (transferable); Council B: revisit at 10k users
F39	Knowledge Graph / Entity Enrichment Entity resolution across extractions; long-term differentiator, premature before scale.	0.96	6	ChatGPT, Perplexity; Diffbot benchmark

ID	Feature & council note	Pri	Effort	Sources
F45	Mobile / PWA Optimisation Desk-work product; lowest council enthusiasm; PWA basics only.	0.86	2	DeepSeek (6.5, lowest), Copilot (transferable)

8. Measurement Plan and Re-scoring Triggers

Instrument first (F50 ships in R1): activation rate (first successful extraction), D1/D7/D30 retention, extractions per session, WAU/MAU, free→paid conversion, organic traffic and indexed pages, referral K-factor, extraction success rate and per-field accuracy.

Benchmarks adopted from the source reports: activation 20–40% is the product-led norm — below 20%, fix onboarding before spending on acquisition; 'good' free→paid is 3–5%, 'great' is 8%+; ungated-experience conversion 7–9%; weekly-habit users convert at 3–4× sporadic users; B2B stickiness norm WAU/MAU 40–60%.

Re-scoring triggers (council-agreed): if analytics show developers/AI-agent builders dominate, pull F21/F23/F24 (API, MCP, LLM markdown) forward into R3 and push F27/F28 back. If sales/lead-gen dominates, pull F28 email verification and F30 CRM forward and delay the developer track. If monitoring beta fails to lift D30 within one quarter, reprioritise integrations (F18/F19) over F32. Pursue SOC 2 (F41) only when an enterprise pipeline justifies it; revisit marketplace (F40) at ~10k MAU.

9. Risks and Caveats

Legal/compliance: public permalinks republish scraped third-party content — ship F04 with a DMCA takedown process and result-cache controls (FD2). Credentialed/paywalled extraction (F33) is excluded from the 18-month plan on CFAA/ToS/GDPR grounds. FD3 (robots.txt compliance + per-domain rate limits) protects shared infrastructure from bans.

Competitive: Firecrawl (developer/agent lane), Browse AI (no-code monitoring) and Clay (enrichment) are all well funded; DatIQ's defensible ground is the no-code, breadth-per-URL, 30-second insight for non-developers — the roadmap deliberately avoids head-on fights (raw-API price war, agent race) until R7.

Evidence limits: several source reports could not see the app behind the SPA and reconstructed features from metadata; the council verified against the live site, but in-app details (export presence, schedule depth, actual conversion data) should be validated by the team against internal analytics before committing R3+ scope. Effort figures are planning-grade estimates, not commitments. Market-size and benchmark figures are third-party aggregates.

Source integrity: Copilot and Meta AI analysed different products bearing similar names; their product-specific recommendations were excluded and only pattern-level ideas retained — worth remembering when commissioning future model research (always confirm the model resolved datiq.app, not datiq.netlify.app, DAT iQ freight, or dating-app namesakes).

10. Bottom Line

Ship R1 immediately: ten small features, ~5.5 person-months, that convert the traffic DatIQ already gets and remember the users it already wins. Launch loudly with R2 once every extraction is a shareable, indexed asset. Make the product a weekly habit in R3 with monitoring — the single feature every model and every council member agrees on — hardened by real infrastructure. Then open the ecosystem (R4), win on being right (R5), sell to teams (R6), and only then chase the agentic frontier (R7) from a position of reliability and distribution strength.